

CS 290B

Scalable Internet Services

Andrew Mutz

October 23, 2014

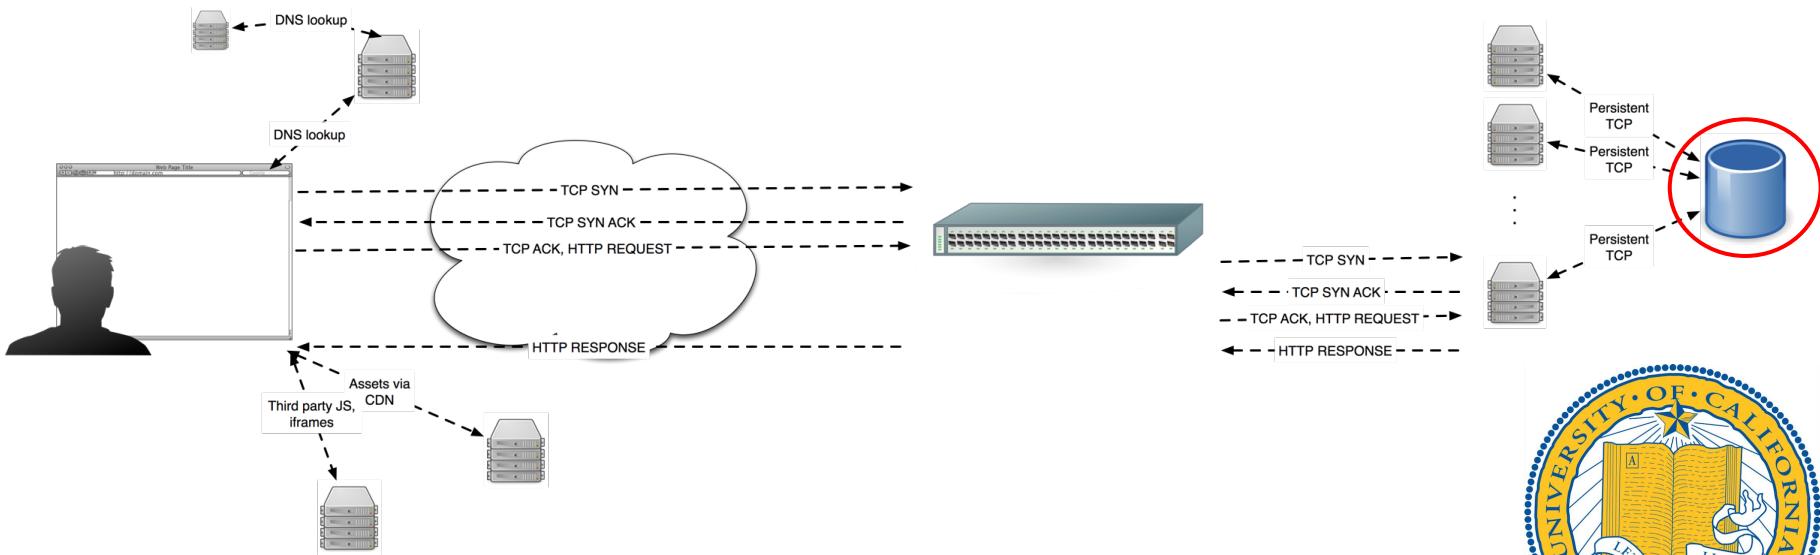


Today's Agenda

A Stable Data Layer: Motivation
Database Concurrency Control
User Authentication with Devise
For Next Time...



Title

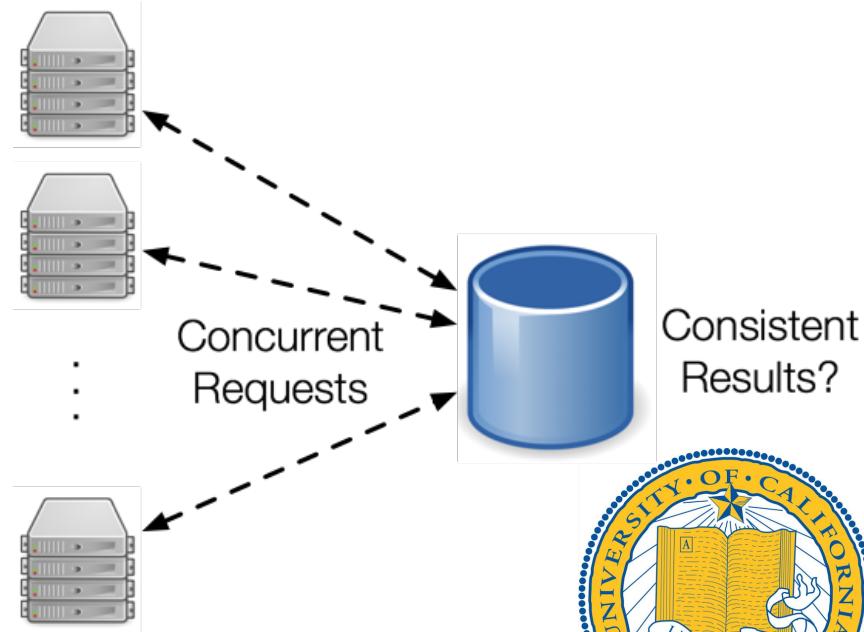


A Stable Data Layer - Transactions

We have many application servers running in parallel.

Each needs to persist data that persists between requests.

The prevailing way to do this today is with Relational Databases

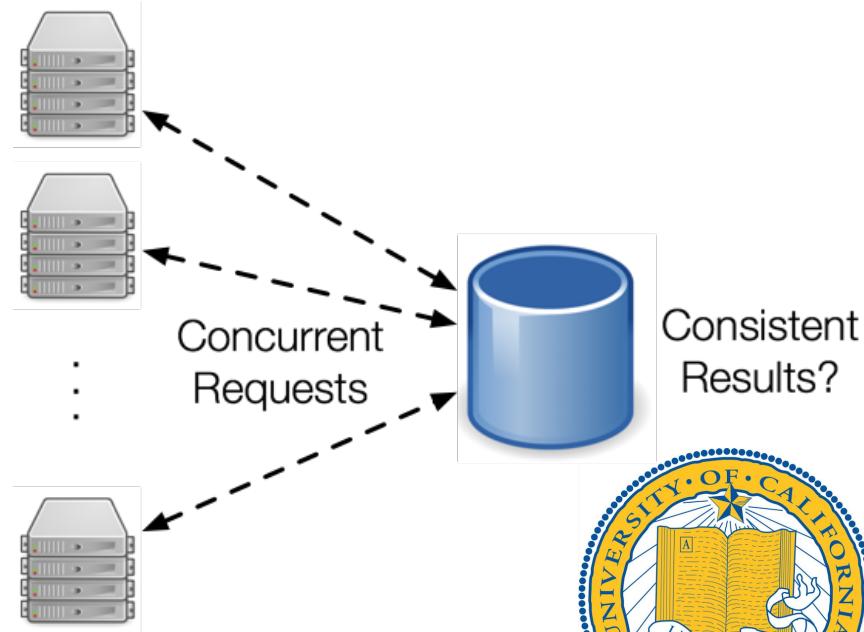


A Stable Data Layer - Transactions

These application servers have needs

- Data needs to be seen by other requests/servers
- Access shouldn't be slow
- Data layer should make sense:

david.withdrawal(100)
mary.deposit(100)



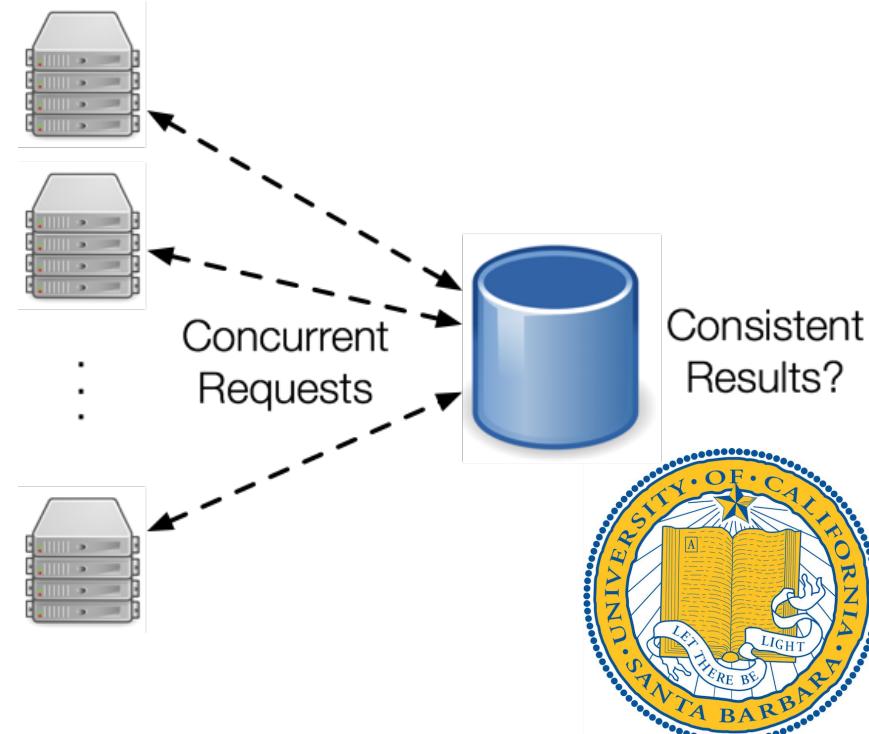
A Stable Data Layer - Transactions

Using transactions in Rails is easy:

```
ActiveRecord::Base.transaction do  
  david.withdrawal(100)  
  mary.deposit(100)  
end
```

Today we will learn more about how these transactions are implemented.

Next time we will learn about transactions in practice



A Stable Data Layer

Database Transactions

- Background
 - Concept that allows a system to guarantee certain semantic properties. Gives control over concurrency.
 - Rigorously defined guarantees mean we can build correct systems on top of them.



A Stable Data Layer - Transactions

History of Database Transactions

- Mid 1970's, IBM System R's RSS
 - First system to implement SQL
 - Introduced formal notions of transactions and serializability
 - Led by Jim Gray
 - Result: 1998 Turing Award.



A Stable Data Layer - Transactions

ACID properties in a database:

- Atomicity
 - All or nothing.
 - No partial application of a transaction.
- Consistency
 - At the beginning and at the end of the transaction, the database should be consistent.
 - Consistency is defined by the integrity constraints



A Stable Data Layer - Transactions

ACID properties in a database:

- Isolation
 - A transaction should not see the effects of other uncommitted transactions.
- Durability
 - Once committed, the transaction's effects should not disappear. (being overwritten by later transactions is fine)



A Stable Data Layer - Transactions

These have overlapping concerns

- Atomicity and Durability are related and are generally provided by journaling
- Consistency and Isolation are provided by concurrency control (usually implemented via locking)

No help with side-effects

- Actions that are visible outside of the system
- Transfer money, communicate with web service, etc.



A Stable Data Layer - Transactions

Schedule (or “history”):

- Abstract model used to describe execution of transactions running in the system.

T1	R(X), W(X), Com.		
T2		R(Y), W(Y), Com.	
T3			R(Z), W(Z), Com.



A Stable Data Layer - Transactions

Conflicting Actions:

- Two actions are said to be in conflict if
 - The actions belong to different transactions
 - At least one of the actions is a write operation
 - The actions access the same object (read or write)
- Example of conflicting actions:
 - T1: R(X), T2: W(X), T3: W(X)
- And these are not conflicting:
 - T1: R(X), T2: R(X), T3: R(X)
 - T1: R(X), T2: W(Y), T3: R(X)

Conflict => we can't blindly execute them in parallel.



A Stable Data Layer - Transactions

Why can't we blindly execute them in parallel? Example:

Lost Update Problem

2nd transaction writes a value on top of a 1st transaction and the value is lost to other transactions running concurrently with the 1st transaction.
Concurrent read transactions will have incorrect results.

T1	R(X)	R(X)	R(X)...
T2		W(X)	Com.



A Stable Data Layer - Transactions

Why can't we blindly execute them in parallel? Example:

Dirty Read Problem

Transactions read a value written by a transaction that is later aborted and removed from the database. Reading transactions will have incorrect results.

T1	R(X)	W(Y), Com.
T2	W(X)	Abort



A Stable Data Layer - Transactions

Why can't we blindly execute them in parallel? Example:

Incorrect Summary Problem

1st transaction takes a summary over the values of all the instances of a repeated data item. While a 2nd transaction updates some instances of the data item. Resulting summary will not reflect a correct result for any deterministic order of the transactions. Result will be random depending on the timing of the updates.

T1	R(All X), AVG	W(Y)	Com.
T2	W(Some X), Com.		



A Stable Data Layer - Transactions

A schedule is **serial** if

- The transactions are executed non-interleaved

Two schedules are **conflict equivalent** if

- They involve the same actions of the same transactions
- Every pair of conflicting actions is ordered in the same way

Schedule S is **conflict serializable** if

- S is conflict equivalent to some serial schedule

A schedule is **recoverable** if

- Transactions commit only after all transactions whose changes they read, commit.



A Stable Data Layer - Transactions

Example of a schedule that is not **conflict serializable**:

T1	R(A), W(A)		R(B), W(B)
T2		R(A), W(A), R(B), W(B)	

Because it is not **conflict equivalent** to this:

T1	R(A), W(A), R(B), W(B)	
T2		R(A), W(A), R(B), W(B)

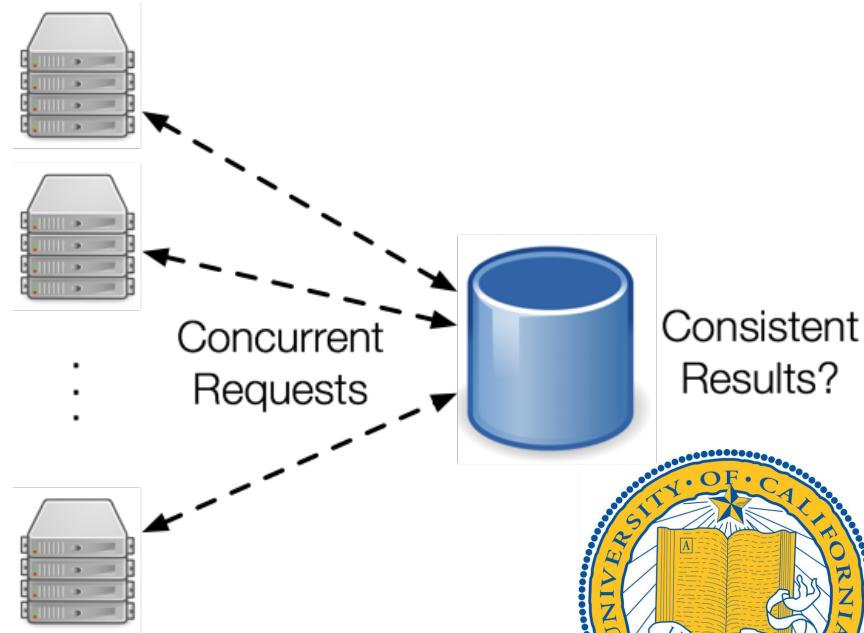
or this:

T1		R(A), W(A), R(B), W(B)
T2	R(A), W(A), R(B), W(B)	



A Stable Data Layer - Transactions

Why is it important that we get a serializable schedule?



A Stable Data Layer - Transactions

Otherwise you can get inconsistent results - not good when you are keeping track of your bank balance in the database

A serial execution of transactions is safe but slow

Most general purpose relational databases default to employing conflict-serializable and recoverable schedules

If you don't want to do a serial execution, what else can you do?

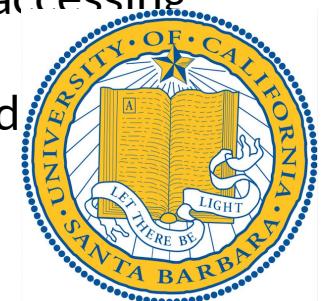


A Stable Data Layer - Transactions

How do we implement a database that schedules that are conflict serializable and recoverable?

Locks

- A lock is a system object associated with a shared resource such as a data item, a row, or a page in memory
- A database lock may need to be acquired by a transaction before accessing the object
- Prevent undesired, incorrect, or inconsistent operations on shared resources by concurrent transactions



A Stable Data Layer - Transactions

Two types of database locks:

- Write-lock
 - Blocks writes and reads
 - Also called “exclusive lock”
- Read-lock
 - Blocks writes
 - Also called “shared lock”



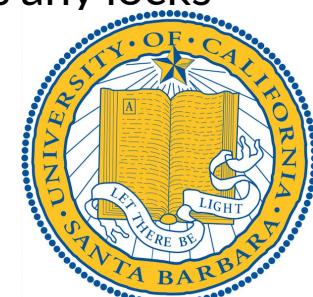
A Stable Data Layer - Transactions

Two-Phase Locking

- 2PL is a concurrency control method that guarantees serializability
- Two-Phase Locking Protocol
 - Each Transaction must obtain a S (shared) lock on object before reading, and an X (exclusive) lock on object before writing
 - If a Transaction holds an X lock on an object, no other Transaction can get a lock (S or X) on that object
 - A transaction cannot request additional locks once it releases any locks
 - Two phases: acquire locks, release locks
- Issue: can result in “cascading aborts”

T1: R(A) W(A) unlock(A) abort

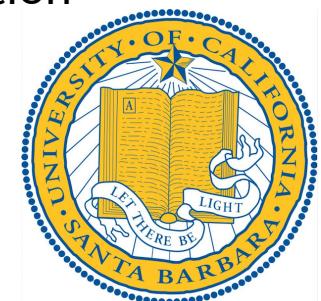
T2: R(A) abort



A Stable Data Layer - Transactions

Strong Strict Two-Phase Locking

- SS2PL allows only conflict serializable schedules
- SS2PL Protocol
 - Each Transaction must obtain a S (shared) lock on object before reading, and an X (exclusive) lock on object before writing
 - If a Transaction holds an X lock on an object, no other Transaction can get a lock (S or X) on that object
 - All locks held by a transaction are released when the transaction completes
- Avoids cascading aborts



Today's Agenda

Up Next: User Authentication with Devise

For Next Time...

- Code, code, code!

